

08-statistical-analysis

September 1, 2026

1 1. Objective

This notebook adds inferential statistics to the objective XAI metrics prepared by 07-metrics-analysis.ipynb. It **loads the existing per-video CSV results** and does not recompute explanations, model predictions, or expensive metrics.

The same videos were evaluated with every XAI method, so method comparisons are repeated-measures comparisons. The Friedman test is the non-parametric omnibus test across all methods; when it is significant, paired Wilcoxon signed-rank tests locate pairwise differences and Holm correction controls family-wise error. These tests preserve the pairing by video and do not require normally distributed observations. Numerical rank alone is not evidence of statistically significant superiority.

1.1 2. Setup and Imports

```
[13]: from itertools import combinations
      from pathlib import Path
      import warnings

      import matplotlib.pyplot as plt
      import numpy as np
      import pandas as pd
      from IPython.display import Markdown, display
      from scipy import stats

      SEED = 42
      ALPHA = 0.05
      np.random.seed(SEED)

      PROJECT_ROOT = Path.cwd()
      if not (PROJECT_ROOT / 'artifacts').exists():
          PROJECT_ROOT = PROJECT_ROOT.parent

      RESULTS_DIR = PROJECT_ROOT / 'docs' / 'results'
      FIGURES_DIR = PROJECT_ROOT / 'docs' / 'figures' / 'statistical-analysis'
      FIGURES_DIR.mkdir(parents=True, exist_ok=True)

      MODEL = 'TANet'
```

```

DATASETS = ['EtriActivity3D', 'Kinetics400']
SCOPES = ['Overall'] + DATASETS
METHODS = [
    '3D-Kernel-SHAP-NEW', '3D-LIME-NEW', '3D-RISE-NEW',
    '3D-Sampled-Occl-Sens-NEW', 'LV-LOCO-NEW', 'LV-Univ-Pred-NEW',
    'SaliencyTubes', 'AOSA', 'GradCAM',
]
DISPLAY_NAMES = {
    '3D-Kernel-SHAP-NEW': '3D-Kernel-SHAP',
    '3D-LIME-NEW': '3D-LIME',
    '3D-RISE-NEW': '3D-RISE',
    '3D-Sampled-Occl-Sens-NEW': '3D-Sampled-Occl-Sens',
    'LV-LOCO-NEW': 'LV-LOCO',
    'LV-Univ-Pred-NEW': 'LV-Univ-Pred',
    'SaliencyTubes': 'SaliencyTubes', 'AOSA': 'AOSA', 'GradCAM': 'GradCAM',
}

# Central direction configuration, reused from notebook 07.
METRIC_CONFIG = {
    'deletion_auc': {'higher_is_better': False, 'family': 'explanation_quality'},
    'insertion_auc': {'higher_is_better': True, 'family': 'explanation_quality'},
    'avg_drop': {'higher_is_better': False, 'family': 'explanation_quality'},
    'pointing_game': {'higher_is_better': True, 'family': 'explanation_quality'},
    'iou50': {'higher_is_better': True, 'family': 'explanation_quality'},
    'iou_value': {'higher_is_better': True, 'family': 'explanation_quality'},
    'bbox_energy_ratio': {'higher_is_better': True, 'family': 'explanation_quality'},
    'temporal_consistency': {'higher_is_better': True, 'family': 'explanation_quality'},
    'sparsity_80': {'higher_is_better': False, 'family': 'explanation_quality'},
    'runtime_seconds': {'higher_is_better': False, 'family': 'computational_cost'},
    'sensitivity_n': {'higher_is_better': True, 'family': 'sensitivity'},
    'roar': {'higher_is_better': True, 'family': 'retraining'},
}

plt.rcParams.update({'font.size': 10, 'axes.grid': True, 'grid.alpha': 0.25})

```

1.2 3. Load Existing Metric Results

Notebook 07 generated three core tidy per-video files and may also provide a Sensitivity-N cache. Original quality metrics use its dataset-local `video_idx`; their pairing therefore inherits notebook 07's assumption that every method's JSON arrays use the same within-dataset video order. Custom metrics, runtime, and Sensitivity-N use the stronger saved class/video-name key. Dataset

is included in every block identifier, so the 60-video overall scope cannot collide across datasets. The runtime convention from notebook 07 is retained (its GradCAM timing source is the same timing file used for SaliencyTubes). ROAR is tracked in the shared taxonomy but is not treated as a repeated-video metric because its cache is aggregated at dataset $\times$ method $\times$ removal-fraction level.

```
[14]: INPUT_FILES = {
    'original': RESULTS_DIR / 'metrics_per_video.csv',
    'custom': RESULTS_DIR / 'custom_metrics_per_video.csv',
    'runtime': RESULTS_DIR / 'runtime_per_video.csv',
}
OPTIONAL_INPUT_FILES = {
    'sensitivity': RESULTS_DIR / 'sensitivity_n_per_video.csv',
    'roar': RESULTS_DIR / 'roar_results.csv',
}
TAXONOMY_PATH = RESULTS_DIR / 'metric_taxonomy.csv'
missing_files = [str(path) for path in INPUT_FILES.values() if not path.
    ↪exists()]
if missing_files:
    raise FileNotFoundError(
        'Run notebooks/07-metrics-analysis.ipynb first. Missing: ' + ', '.
    ↪join(missing_files)
    )

source_frames = {name: pd.read_csv(path) for name, path in INPUT_FILES.items()}

# Sensitivity-N is an optional per-video cache. Apply notebook 07's canonical
    ↪configuration.
SENSITIVITY_CONFIG = {'n': 10, 'n_subsets': 100, 'seed': 777}
sensitivity_path = OPTIONAL_INPUT_FILES['sensitivity']
if sensitivity_path.exists():
    sensitivity_frame = pd.read_csv(sensitivity_path)
    if not sensitivity_frame.empty:
        for column, expected in SENSITIVITY_CONFIG.items():
            sensitivity_frame = sensitivity_frame[sensitivity_frame[column].
    ↪eq(expected)]
        sensitivity_frame = sensitivity_frame[sensitivity_frame['model'].
    ↪eq(MODEL)].copy()
        sensitivity_frame['method_display'] = sensitivity_frame['method'].
    ↪map(DISPLAY_NAMES)
        source_frames['sensitivity'] = sensitivity_frame

taxonomy_df = (pd.read_csv(TAXONOMY_PATH) if TAXONOMY_PATH.exists()
    else pd.DataFrame({'metric': list(METRIC_CONFIG),
        'category': [METRIC_CONFIG[m]['family']]
    ↪for m in METRIC_CONFIG}))
taxonomy_metrics = set(taxonomy_df['metric'])
```

```
{name: frame.shape for name, frame in source_frames.items()}
```

```
[14]: {'original': (540, 11),
      'custom': (540, 9),
      'runtime': (540, 7),
      'sensitivity': (540, 12)}
```

```
[15]: METADATA_COLUMNS = {
      'dataset', 'model', 'method', 'method_display',
      'video_idx', 'video_class', 'video_name',
    }

    def discover_metric_columns(frame):
        # Notebook 07's taxonomy is authoritative, so diagnostic/control columns
        ↪are not misclassified.
        return [column for column in frame.columns if column in taxonomy_metrics]

    def to_metric_long(frame, source_name):
        metric_columns = discover_metric_columns(frame)
        if 'video_idx' in frame.columns:
            sample_id = frame['dataset'].astype(str) + '::index::' +
            ↪frame['video_idx'].astype(str)
        else:
            sample_id = (
                frame['dataset'].astype(str) + '::' + frame['video_class'].
            ↪astype(str)
                + '::' + frame['video_name'].astype(str)
            )
        prepared = frame.assign(sample_id=sample_id, source=source_name)
        return prepared.melt(
            id_vars=['dataset', 'model', 'method', 'method_display', 'sample_id',
            ↪'source'],
            value_vars=metric_columns, var_name='metric', value_name='value',
        ), metric_columns

    long_parts = []
    discovered_by_source = {}
    for source_name, frame in source_frames.items():
        part, columns = to_metric_long(frame, source_name)
        long_parts.append(part)
        discovered_by_source[source_name] = columns

    metric_long = pd.concat(long_parts, ignore_index=True)
    metric_long['value'] = pd.to_numeric(metric_long['value'], errors='coerce')
    METRICS = list(dict.fromkeys(metric for columns in discovered_by_source.
    ↪values() for metric in columns))
    unconfigured = sorted(set(METRICS) - set(METRIC_CONFIG))
```

```

if unconfigured:
    raise ValueError(f'Add metric direction(s) to METRIC_CONFIG:␣
↳{unconfigured}')
taxonomy_category = taxonomy_df.set_index('metric')['category'].to_dict()
metric_long['family'] = metric_long['metric'].map(taxonomy_category).fillna(
    metric_long['metric'].map(lambda metric: METRIC_CONFIG[metric]['family']))
)

availability_rows = []
for metric in taxonomy_df['metric']:
    observed = metric_long[metric_long['metric'].eq(metric)]
    if metric == 'roar':
        reason = 'Excluded: results are dataset/method/fraction-level, not␣
↳per-video repeated measures.'
        included = False
    elif observed.empty:
        reason = 'Unavailable: no per-video cache rows.'
        included = False
    else:
        reason = 'Included; complete cases are determined per scope below.'
        included = True
    availability_rows.append({
        'Metric': metric, 'Category': taxonomy_category.get(metric),
        'Valid per-video values': int(observed['value'].notna().sum()),
        'Methods observed': int(observed['method'].nunique()),
        'Included in repeated-video analysis': included, 'Reason': reason,
    })
metric_availability_df = pd.DataFrame(availability_rows)
metric_availability_df.to_csv(RESULTS_DIR / 'statistical_metric_availability.
↳csv', index=False)

display(pd.DataFrame([
    {
        'Metric': metric,
        'Source': next(source for source, columns in discovered_by_source.
↳items() if metric in columns),
        'Direction': 'higher is better' if␣
↳METRIC_CONFIG[metric]['higher_is_better'] else 'lower is better',
        'Family': taxonomy_category.get(metric,␣
↳METRIC_CONFIG[metric]['family']),
    }
    for metric in METRICS
]))
display(metric_availability_df)

```

	Metric	Source	Direction \
0	deletion_auc	original	lower is better

1	insertion_auc	original	higher is better
2	avg_drop	original	lower is better
3	pointing_game	original	higher is better
4	iou50	original	higher is better
5	iou_value	original	higher is better
6	bbox_energy_ratio	custom	higher is better
7	temporal_consistency	custom	higher is better
8	sparsity_80	custom	lower is better
9	runtime_seconds	runtime	lower is better
10	sensitivity_n	sensitivity	higher is better

	Family
0	input_perturbation
1	input_perturbation
2	input_perturbation
3	ground_truth_comparison
4	ground_truth_comparison
5	ground_truth_comparison
6	project_specific
7	project_specific
8	project_specific
9	computational_cost
10	sensitivity

	Metric	Category	Valid per-video values \
0	deletion_auc	input_perturbation	540
1	insertion_auc	input_perturbation	540
2	avg_drop	input_perturbation	540
3	pointing_game	ground_truth_comparison	540
4	iou50	ground_truth_comparison	540
5	iou_value	ground_truth_comparison	540
6	sensitivity_n	sensitivity	540
7	roar	retraining	0
8	bbox_energy_ratio	project_specific	540
9	temporal_consistency	project_specific	540
10	sparsity_80	project_specific	540
11	runtime_seconds	computational_cost	540

	Methods observed	Included in repeated-video analysis \
0	9	True
1	9	True
2	9	True
3	9	True
4	9	True
5	9	True
6	9	True
7	0	False
8	9	True

9	9	True
10	9	True
11	9	True

	Reason
0	Included; complete cases are determined per sc...
1	Included; complete cases are determined per sc...
2	Included; complete cases are determined per sc...
3	Included; complete cases are determined per sc...
4	Included; complete cases are determined per sc...
5	Included; complete cases are determined per sc...
6	Included; complete cases are determined per sc...
7	Excluded: results are dataset/method/fraction-...
8	Included; complete cases are determined per sc...
9	Included; complete cases are determined per sc...
10	Included; complete cases are determined per sc...
11	Included; complete cases are determined per sc...

1.3 4. Data Validation

Every statistical block must contain at most one value per dataset, metric, method, and sample. Coverage is checked against the nine configured methods. Complete-case filtering happens separately for each metric and scope; the global-test table reports both available blocks and rows removed.

```
[16]: if set(metric_long['model'].dropna().unique()) != {MODEL}:
        raise ValueError(f'Unexpected models: {sorted(metric_long["model"].dropna().
        ↪unique())}')
    if set(metric_long['dataset'].unique()) != set(DATASETS):
        raise ValueError(f'Unexpected datasets: {sorted(metric_long["dataset"].
        ↪unique())}')
    if set(metric_long['method'].unique()) != set(METHODS):
        raise ValueError('Observed XAI methods do not match the nine methods_
        ↪configured in notebook 07.')

    key_columns = ['dataset', 'metric', 'method', 'sample_id']
    duplicate_count = int(metric_long.duplicated(key_columns).sum())
    if duplicate_count:
        raise ValueError(f'{duplicate_count} duplicate repeated-measure keys found.
        ↪')

    validation_df = (
        metric_long.groupby(['source', 'dataset', 'metric'])
        .agg(rows=('value', 'size'), valid_values=('value', 'count'),
            videos=('sample_id', 'nunique'), methods=('method', 'nunique'))
        .reset_index()
    )
```

```

validation_df['missing_values'] = validation_df['rows'] -
    ↪validation_df['valid_values']
validation_df['expected_rows'] = validation_df['videos'] * len(METHODS)
validation_df['complete_shape'] = (
    (validation_df['rows'] == validation_df['expected_rows'])
    & (validation_df['methods'] == len(METHODS))
)
display(validation_df)
print(f'Duplicate keys: {duplicate_count}; missing values: {int(metric_long.
    ↪value.isna().sum())}')

```

	source	dataset	metric	rows	valid_values \
0	custom	EtriActivity3D	bbox_energy_ratio	270	270
1	custom	EtriActivity3D	sparsity_80	270	270
2	custom	EtriActivity3D	temporal_consistency	270	270
3	custom	Kinetics400	bbox_energy_ratio	270	270
4	custom	Kinetics400	sparsity_80	270	270
5	custom	Kinetics400	temporal_consistency	270	270
6	original	EtriActivity3D	avg_drop	270	270
7	original	EtriActivity3D	deletion_auc	270	270
8	original	EtriActivity3D	insertion_auc	270	270
9	original	EtriActivity3D	iou50	270	270
10	original	EtriActivity3D	iou_value	270	270
11	original	EtriActivity3D	pointing_game	270	270
12	original	Kinetics400	avg_drop	270	270
13	original	Kinetics400	deletion_auc	270	270
14	original	Kinetics400	insertion_auc	270	270
15	original	Kinetics400	iou50	270	270
16	original	Kinetics400	iou_value	270	270
17	original	Kinetics400	pointing_game	270	270
18	runtime	EtriActivity3D	runtime_seconds	270	270
19	runtime	Kinetics400	runtime_seconds	270	270
20	sensitivity	EtriActivity3D	sensitivity_n	270	270
21	sensitivity	Kinetics400	sensitivity_n	270	270

	videos	methods	missing_values	expected_rows	complete_shape
0	30	9	0	270	True
1	30	9	0	270	True
2	30	9	0	270	True
3	30	9	0	270	True
4	30	9	0	270	True
5	30	9	0	270	True
6	30	9	0	270	True
7	30	9	0	270	True
8	30	9	0	270	True
9	30	9	0	270	True
10	30	9	0	270	True
11	30	9	0	270	True

12	30	9	0	270	True
13	30	9	0	270	True
14	30	9	0	270	True
15	30	9	0	270	True
16	30	9	0	270	True
17	30	9	0	270	True
18	30	9	0	270	True
19	30	9	0	270	True
20	30	9	0	270	True
21	30	9	0	270	True

Duplicate keys: 0; missing values: 0

1.4 5. Descriptive Statistics

The table below contains N, mean, sample standard deviation, median, quartiles, and IQR for every dataset $\times$ method $\times$ discovered metric, plus the combined 60-video scope. It is saved in long form for direct thesis use.

```
[17]: def descriptive_statistics(data):
    scoped = pd.concat([
        data.assign(Dataset=data['dataset']),
        data.assign(Dataset='Overall'),
    ], ignore_index=True)
    summary = (
        scoped.groupby(['metric', 'Dataset', 'method',
            ↪ 'method_display'])['value']
        .agg(
            N='count', mean='mean', std='std', median='median',
            q25=lambda values: values.quantile(0.25),
            q75=lambda values: values.quantile(0.75),
        )
        .reset_index()
    )
    summary['IQR'] = summary['q75'] - summary['q25']
    summary['direction'] = summary['metric'].map(
        lambda metric: 'higher' if METRIC_CONFIG[metric]['higher_is_better']
        ↪ else 'lower'
    )
    return summary.sort_values(['metric', 'Dataset', 'method_display']).
    ↪ reset_index(drop=True)

descriptive_df = descriptive_statistics(metric_long)
descriptive_df.to_csv(RESULTS_DIR / 'statistical_descriptive_summary.csv',
    ↪ index=False)
display(descriptive_df.round(4))
```

	metric	Dataset	method \
0	avg_drop	EtriActivity3D	3D-Kernel-SHAP-NEW

```

1          avg_drop EtriActivity3D          3D-LIME-NEW
2          avg_drop EtriActivity3D          3D-RISE-NEW
3          avg_drop EtriActivity3D 3D-Sampled-Occl-Sens-NEW
4          avg_drop EtriActivity3D          AOSA
..          ...
292 temporal_consistency Overall          AOSA
293 temporal_consistency Overall          GradCAM
294 temporal_consistency Overall          LV-LOCO-NEW
295 temporal_consistency Overall          LV-Univ-Pred-NEW
296 temporal_consistency Overall          SaliencyTubes

          method_display  N    mean    std  median    q25    q75  \
0          3D-Kernel-SHAP 30  22.2154  24.6748  17.1348  0.0009  41.3132
1          3D-LIME 30  16.3615  26.3627   0.6543  0.0000  26.9995
2          3D-RISE 30   2.3191   7.1667   0.0001  0.0000   0.9070
3  3D-Sampled-Occl-Sens 30  53.4685  44.1555  72.4849  4.3768  97.5397
4          AOSA 30  20.9533  31.9296   0.4261  0.0000  39.0462
..          ... ..
292          AOSA 60   0.9312   0.0599   0.9504  0.8871   0.9827
293          GradCAM 60   0.9375   0.0463   0.9466  0.9097   0.9694
294          LV-LOCO 60   0.9227   0.0479   0.9258  0.8842   0.9694
295          LV-Univ-Pred 60   0.9477   0.0432   0.9539  0.9153   0.9884
296          SaliencyTubes 60   0.9741   0.0238   0.9814  0.9672   0.9932

          IQR direction
0  41.3123  lower
1  26.9995  lower
2   0.9070  lower
3  93.1629  lower
4  39.0462  lower
..          ...
292  0.0956  higher
293  0.0598  higher
294  0.0852  higher
295  0.0731  higher
296  0.0259  higher

```

[297 rows x 12 columns]

1.5 6. Friedman Tests

For each scope and metric, columns are the nine methods and rows are video blocks. A row enters the test only when all nine values are finite. The null hypothesis is that the XAI methods have equivalent distributions/performance for that metric. Kendall's W is reported as `chi-square / (N * (k - 1))`, using the tie-corrected Friedman statistic returned by SciPy.

```
[18]: def complete_case_matrix(data, metric, scope):
      subset = data[data['metric'].eq(metric)]
```

```

    if scope != 'Overall':
        subset = subset[subset['dataset'].eq(scope)]
        matrix = subset.pivot(index='sample_id', columns='method', values='value').
        ↪reindex(columns=METHODS)
        matrix = matrix.replace([np.inf, -np.inf], np.nan)
        return matrix, matrix.dropna(axis=0, how='any')

def friedman_test(data, metric, scope):
    available, complete = complete_case_matrix(data, metric, scope)
    n, k = complete.shape
    if n < 2 or k < 3:
        statistic, p_value = np.nan, np.nan
    elif np.all(np.ptp(complete.to_numpy(dtype=float), axis=1) == 0):
        statistic, p_value = 0.0, 1.0
    else:
        statistic, p_value = stats.friedmanchisquare(
            *[complete[method].to_numpy() for method in METHODS]
        )
    kendall_w = statistic / (n * (k - 1)) if n and np.isfinite(statistic) else ↪
    ↪np.nan
    return {
        'Metric': metric, 'Dataset': scope, 'N': n, 'Methods': k,
        'Available blocks': len(available), 'Rows removed': len(available) - n,
        'Friedman statistic': statistic, 'p': p_value, 'Kendall W': kendall_w,
        'Significant': bool(np.isfinite(p_value) and p_value < ALPHA),
    }

global_tests_df = pd.DataFrame([
    friedman_test(metric_long, metric, scope)
    for metric in METRICS for scope in SCOPES
])
global_tests_df.to_csv(RESULTS_DIR / 'friedman_global_tests.csv', index=False)
display(global_tests_df.round({'Friedman statistic': 4, 'p': 6, 'Kendall W': ↪
    ↪4}))

```

	Metric	Dataset	N	Methods	Available blocks \
0	deletion_auc	Overall	60	9	60
1	deletion_auc	EtriActivity3D	30	9	30
2	deletion_auc	Kinetics400	30	9	30
3	insertion_auc	Overall	60	9	60
4	insertion_auc	EtriActivity3D	30	9	30
5	insertion_auc	Kinetics400	30	9	30
6	avg_drop	Overall	60	9	60
7	avg_drop	EtriActivity3D	30	9	30
8	avg_drop	Kinetics400	30	9	30
9	pointing_game	Overall	60	9	60
10	pointing_game	EtriActivity3D	30	9	30

11	pointing_game	Kinetics400	30	9	30
12	iou50	Overall	60	9	60
13	iou50	EtriActivity3D	30	9	30
14	iou50	Kinetics400	30	9	30
15	iou_value	Overall	60	9	60
16	iou_value	EtriActivity3D	30	9	30
17	iou_value	Kinetics400	30	9	30
18	bbox_energy_ratio	Overall	60	9	60
19	bbox_energy_ratio	EtriActivity3D	30	9	30
20	bbox_energy_ratio	Kinetics400	30	9	30
21	temporal_consistency	Overall	60	9	60
22	temporal_consistency	EtriActivity3D	30	9	30
23	temporal_consistency	Kinetics400	30	9	30
24	sparsity_80	Overall	60	9	60
25	sparsity_80	EtriActivity3D	30	9	30
26	sparsity_80	Kinetics400	30	9	30
27	runtime_seconds	Overall	60	9	60
28	runtime_seconds	EtriActivity3D	30	9	30
29	runtime_seconds	Kinetics400	30	9	30
30	sensitivity_n	Overall	60	9	60
31	sensitivity_n	EtriActivity3D	30	9	30
32	sensitivity_n	Kinetics400	30	9	30

	Rows removed	Friedman statistic	p	Kendall W	Significant
0	0	136.4356	0.000000	0.2842	True
1	0	83.4400	0.000000	0.3477	True
2	0	69.0756	0.000000	0.2878	True
3	0	115.5067	0.000000	0.2406	True
4	0	108.5600	0.000000	0.4523	True
5	0	79.1022	0.000000	0.3296	True
6	0	187.9001	0.000000	0.3915	True
7	0	94.2963	0.000000	0.3929	True
8	0	133.8724	0.000000	0.5578	True
9	0	35.4225	0.000022	0.0738	True
10	0	26.0503	0.001030	0.1085	True
11	0	18.9257	0.015262	0.0789	True
12	0	25.7273	0.001169	0.0536	True
13	0	24.7000	0.001747	0.1029	True
14	0	9.0000	0.342296	0.0375	False
15	0	42.1493	0.000001	0.0878	True
16	0	54.8266	0.000000	0.2284	True
17	0	18.3361	0.018843	0.0764	True
18	0	149.2773	0.000000	0.3110	True
19	0	115.2000	0.000000	0.4800	True
20	0	53.7471	0.000000	0.2239	True
21	0	324.6133	0.000000	0.6763	True
22	0	170.6756	0.000000	0.7111	True
23	0	180.7111	0.000000	0.7530	True

24	0	194.6444	0.000000	0.4055	True
25	0	96.0444	0.000000	0.4002	True
26	0	129.1111	0.000000	0.5380	True
27	0	475.0073	0.000000	0.9896	True
28	0	238.2342	0.000000	0.9926	True
29	0	236.8359	0.000000	0.9868	True
30	0	129.9156	0.000000	0.2707	True
31	0	82.5956	0.000000	0.3441	True
32	0	68.4800	0.000000	0.2853	True

1.6 7. Post-hoc Wilcoxon Tests

All 36 method pairs are tested only after a significant Friedman result. Each family is one metric $\times$ scope. Holm-adjusted p-values are computed with the step-down formula. All-zero paired differences are defined transparently as statistic 0, $p = 1$, and effect size 0 instead of allowing the test to fail.

```
[19]: def holm_adjust(p_values):
    p_values = np.asarray(p_values, dtype=float)
    order = np.argsort(p_values)
    sorted_p = p_values[order]
    adjusted_sorted = np.minimum(1.0, np.maximum.accumulate(
        (len(sorted_p) - np.arange(len(sorted_p))) * sorted_p
    ))
    adjusted = np.empty_like(adjusted_sorted)
    adjusted[order] = adjusted_sorted
    return adjusted

def paired_rank_biserial(differences):
    differences = np.asarray(differences, dtype=float)
    nonzero = differences[~np.isclose(differences, 0.0)]
    if len(nonzero) == 0:
        return 0.0
    ranks = stats.rankdata(np.abs(nonzero), method='average')
    return float((ranks[nonzero > 0].sum() - ranks[nonzero < 0].sum()) / ranks.
↪sum())

def safe_wilcoxon(a, b):
    differences = np.asarray(a, dtype=float) - np.asarray(b, dtype=float)
    if np.all(np.isclose(differences, 0.0)):
        return 0.0, 1.0, 0.0
    with warnings.catch_warnings():
        warnings.simplefilter('ignore', UserWarning)
        result = stats.wilcoxon(
            a, b, zero_method='wilcox', correction=False,
            alternative='two-sided', method='auto',
        )
```

```

        return float(result.statistic), float(result.pvalue),
        paired_rank_biserial(differences)

def posthoc_for_family(data, metric, scope):
    _, complete = complete_case_matrix(data, metric, scope)
    rows = []
    for method_a, method_b in combinations(METHODS, 2):
        a, b = complete[method_a].to_numpy(), complete[method_b].to_numpy()
        statistic, p_raw, effect = safe_wilcoxon(a, b)
        rows.append({
            'Metric': metric, 'Dataset': scope, 'N': len(complete),
            'Method A': DISPLAY_NAMES[method_a], 'Method B':
            DISPLAY_NAMES[method_b],
            'Wilcoxon statistic': statistic, 'p raw': p_raw,
            'Mean difference (A-B)': float(np.mean(a - b)),
            'Median difference (A-B)': float(np.median(a - b)),
            'Rank-biserial effect': effect,
        })
    family = pd.DataFrame(rows)
    family['p Holm'] = holm_adjust(family['p raw'])
    family['Significant'] = family['p Holm'] < ALPHA
    return family

posthoc_families = []
for row in global_tests_df.itertuples(index=False):
    if row.Significant:
        posthoc_families.append(posthoc_for_family(metric_long, row.Metric, row.
        Dataset))

POSTHOC_COLUMNS = [
    'Metric', 'Dataset', 'N', 'Method A', 'Method B', 'Wilcoxon statistic',
    'p raw', 'p Holm', 'Mean difference (A-B)', 'Median difference (A-B)',
    'Rank-biserial effect', 'Significant',
]
posthoc_df = (pd.concat(posthoc_families, ignore_index=True)[POSTHOC_COLUMNS]
               if posthoc_families else pd.DataFrame(columns=POSTHOC_COLUMNS))
posthoc_df.to_csv(RESULTS_DIR / 'wilcoxon_holm_posthoc.csv', index=False)
display(posthoc_df.round(5))
display(posthoc_df.groupby(['Metric', 'Dataset']).size().rename('Comparisons').
        reset_index())

```

	Metric	Dataset	N	Method A	Method B \
0	deletion_auc	Overall	60	3D-Kernel-SHAP	3D-LIME
1	deletion_auc	Overall	60	3D-Kernel-SHAP	3D-RISE
2	deletion_auc	Overall	60	3D-Kernel-SHAP	3D-Sampled-Occl-Sens
3	deletion_auc	Overall	60	3D-Kernel-SHAP	LV-LOCO
4	deletion_auc	Overall	60	3D-Kernel-SHAP	LV-Univ-Pred

...	...	...	..	...	...
1147	sensitivity_n	Kinetics400	30	LV-Univ-Pred	AOSA
1148	sensitivity_n	Kinetics400	30	LV-Univ-Pred	GradCAM
1149	sensitivity_n	Kinetics400	30	SaliencyTubes	AOSA
1150	sensitivity_n	Kinetics400	30	SaliencyTubes	GradCAM
1151	sensitivity_n	Kinetics400	30	AOSA	GradCAM

	Wilcoxon statistic	p raw	p Holm	Mean difference (A-B)	\
0	317.0	0.00001	0.00025	0.03687	
1	842.0	0.59099	1.00000	0.00520	
2	259.0	0.00000	0.00004	-0.13272	
3	303.0	0.00001	0.00017	-0.08273	
4	137.0	0.00000	0.00000	-0.23050	
...	...	...	...	...	
1147	74.0	0.00067	0.01600	-0.18762	
1148	27.0	0.00000	0.00008	-0.28060	
1149	181.0	0.29884	1.00000	0.03426	
1150	71.0	0.00051	0.01314	-0.05872	
1151	113.0	0.01283	0.24384	-0.09298	

	Median difference (A-B)	Rank-biserial effect	Significant
0	0.02346	0.65355	True
1	0.00140	0.07978	False
2	-0.08973	-0.71694	True
3	-0.05620	-0.66885	True
4	-0.24246	-0.85027	True
...	...	...	...
1147	-0.20164	-0.68172	True
1148	-0.30159	-0.88387	True
1149	0.06147	0.22151	False
1150	-0.04794	-0.69462	True
1151	-0.08448	-0.51398	False

[1152 rows x 12 columns]

	Metric	Dataset	Comparisons
0	avg_drop	EtriActivity3D	36
1	avg_drop	Kinetics400	36
2	avg_drop	Overall	36
3	bbox_energy_ratio	EtriActivity3D	36
4	bbox_energy_ratio	Kinetics400	36
5	bbox_energy_ratio	Overall	36
6	deletion_auc	EtriActivity3D	36
7	deletion_auc	Kinetics400	36
8	deletion_auc	Overall	36
9	insertion_auc	EtriActivity3D	36
10	insertion_auc	Kinetics400	36
11	insertion_auc	Overall	36

12	iou50	EtriActivity3D	36
13	iou50	Overall	36
14	iou_value	EtriActivity3D	36
15	iou_value	Kinetics400	36
16	iou_value	Overall	36
17	pointing_game	EtriActivity3D	36
18	pointing_game	Kinetics400	36
19	pointing_game	Overall	36
20	runtime_seconds	EtriActivity3D	36
21	runtime_seconds	Kinetics400	36
22	runtime_seconds	Overall	36
23	sensitivity_n	EtriActivity3D	36
24	sensitivity_n	Kinetics400	36
25	sensitivity_n	Overall	36
26	sparsity_80	EtriActivity3D	36
27	sparsity_80	Kinetics400	36
28	sparsity_80	Overall	36
29	temporal_consistency	EtriActivity3D	36
30	temporal_consistency	Kinetics400	36
31	temporal_consistency	Overall	36

1.7 8. Effect Sizes

Kendall's W accompanies each Friedman test. For paired Wilcoxon comparisons, rank-biserial correlation is $(W+ - W-) / (W+ + W-)$ after zero differences are removed; its sign describes A minus B on the metric's raw scale. Magnitude, practical importance, and the configured metric direction should be interpreted together.

```
[20]: effect_size_summary = global_tests_df[[
    'Metric', 'Dataset', 'N', 'Kendall W', 'Significant'
]].copy()
display(effect_size_summary.round(4))
if not posthoc_df.empty:
    display(posthoc_df.loc[posthoc_df['Significant'], [
        'Metric', 'Dataset', 'Method A', 'Method B', 'p Holm', 'Rank-biserial_
↪effect'
    ]].round(4))
```

	Metric	Dataset	N	Kendall W	Significant
0	deletion_auc	Overall	60	0.2842	True
1	deletion_auc	EtriActivity3D	30	0.3477	True
2	deletion_auc	Kinetics400	30	0.2878	True
3	insertion_auc	Overall	60	0.2406	True
4	insertion_auc	EtriActivity3D	30	0.4523	True
5	insertion_auc	Kinetics400	30	0.3296	True
6	avg_drop	Overall	60	0.3915	True
7	avg_drop	EtriActivity3D	30	0.3929	True
8	avg_drop	Kinetics400	30	0.5578	True

9	pointing_game	Overall	60	0.0738	True
10	pointing_game	EtriActivity3D	30	0.1085	True
11	pointing_game	Kinetics400	30	0.0789	True
12	iou50	Overall	60	0.0536	True
13	iou50	EtriActivity3D	30	0.1029	True
14	iou50	Kinetics400	30	0.0375	False
15	iou_value	Overall	60	0.0878	True
16	iou_value	EtriActivity3D	30	0.2284	True
17	iou_value	Kinetics400	30	0.0764	True
18	bbox_energy_ratio	Overall	60	0.3110	True
19	bbox_energy_ratio	EtriActivity3D	30	0.4800	True
20	bbox_energy_ratio	Kinetics400	30	0.2239	True
21	temporal_consistency	Overall	60	0.6763	True
22	temporal_consistency	EtriActivity3D	30	0.7111	True
23	temporal_consistency	Kinetics400	30	0.7530	True
24	sparsity_80	Overall	60	0.4055	True
25	sparsity_80	EtriActivity3D	30	0.4002	True
26	sparsity_80	Kinetics400	30	0.5380	True
27	runtime_seconds	Overall	60	0.9896	True
28	runtime_seconds	EtriActivity3D	30	0.9926	True
29	runtime_seconds	Kinetics400	30	0.9868	True
30	sensitivity_n	Overall	60	0.2707	True
31	sensitivity_n	EtriActivity3D	30	0.3441	True
32	sensitivity_n	Kinetics400	30	0.2853	True

	Metric	Dataset	Method A	Method B \
0	deletion_auc	Overall	3D-Kernel-SHAP	3D-LIME
2	deletion_auc	Overall	3D-Kernel-SHAP	3D-Sampled-Occl-Sens
3	deletion_auc	Overall	3D-Kernel-SHAP	LV-LOCO
4	deletion_auc	Overall	3D-Kernel-SHAP	LV-Univ-Pred
6	deletion_auc	Overall	3D-Kernel-SHAP	AOSA
...	...	...	...	...
1144	sensitivity_n	Kinetics400	LV-LOCO	AOSA
1146	sensitivity_n	Kinetics400	LV-Univ-Pred	SaliencyTubes
1147	sensitivity_n	Kinetics400	LV-Univ-Pred	AOSA
1148	sensitivity_n	Kinetics400	LV-Univ-Pred	GradCAM
1150	sensitivity_n	Kinetics400	SaliencyTubes	GradCAM

	p Holm	Rank-biserial effect
0	0.0002	0.6536
2	0.0000	-0.7169
3	0.0002	-0.6689
4	0.0000	-0.8503
6	0.0087	-0.5180
...	...	...
1144	0.0283	0.6516
1146	0.0053	-0.7376
1147	0.0160	-0.6817

```
1148  0.0001          -0.8839
1150  0.0131          -0.6946
```

[565 rows x 6 columns]

1.8 9. XAI Method Rankings

Methods are ranked by mean performance within each metric and scope, respecting the direction configuration. Mean and median are both retained. Ties receive average ranks. **best** and **worst** identify numerical endpoints only; statistically supported pairwise differences are exclusively those marked significant in the Holm table.

```
[21]: def build_rankings(descriptive):
    ranked_groups = []
    for (metric, scope), group in descriptive.groupby(['metric', 'Dataset'],
    ↪sort=False):
        group = group.copy()
        ascending = not METRIC_CONFIG[metric]['higher_is_better']
        group['Rank'] = group['mean'].rank(method='average',
    ↪ascending=ascending)
        group['Best'] = group['Rank'].eq(group['Rank'].min())
        group['Worst'] = group['Rank'].eq(group['Rank'].max())
        ranked_groups.append(group)
    ranking = pd.concat(ranked_groups, ignore_index=True)
    ranking = ranking.rename(columns={
        'metric': 'Metric', 'method_display': 'Method', 'mean': 'Mean',
    ↪'median': 'Median'
    })
    return ranking.sort_values(['Metric', 'Dataset', 'Rank', 'Method']).
    ↪reset_index(drop=True)

ranking_df = build_rankings(descriptive_df)
ranking_columns = ['Metric', 'Dataset', 'Rank', 'Method', 'Mean', 'Median',
    ↪'N', 'Best', 'Worst']
ranking_df[ranking_columns].to_csv(RESULTS_DIR / 'xai_method_rankings.csv',
    ↪index=False)
display(ranking_df[ranking_columns].round(4))
```

	Metric	Dataset	Rank	Method \
0	avg_drop	EtriActivity3D	1.0	3D-RISE
1	avg_drop	EtriActivity3D	2.0	3D-LIME
2	avg_drop	EtriActivity3D	3.0	SaliencyTubes
3	avg_drop	EtriActivity3D	4.0	AOSA
4	avg_drop	EtriActivity3D	5.0	3D-Kernel-SHAP
..	...	...	...	...
292	temporal_consistency	Overall	5.0	3D-Kernel-SHAP
293	temporal_consistency	Overall	6.0	AOSA
294	temporal_consistency	Overall	7.0	LV-LOCO

295	temporal_consistency	Overall	8.0	3D-RISE
296	temporal_consistency	Overall	9.0	3D-Sampled-Occl-Sens

	Mean	Median	N	Best	Worst
0	2.3191	0.0001	30	True	False
1	16.3615	0.6543	30	False	False
2	20.1009	4.8027	30	False	False
3	20.9533	0.4261	30	False	False
4	22.2154	17.1348	30	False	False
..	...	...	..	...	...
292	0.9350	0.9512	60	False	False
293	0.9312	0.9504	60	False	False
294	0.9227	0.9258	60	False	False
295	0.4457	0.4449	60	False	False
296	0.0080	-0.0137	60	False	True

[297 rows x 9 columns]

1.9 10. Analysis by Dataset

The dataset-specific Friedman rows above compare methods within ETRI and within Kinetics using paired videos. This secondary analysis asks a different question: whether a given method's values differ between datasets. ETRI and Kinetics contain unrelated videos, so a paired test would be invalid. Two-sided Mann-Whitney U tests therefore treat datasets as independent samples; Holm correction is applied across the nine methods separately for each metric. The independent rank-biserial effect is $2U/(n_{\text{ETRI}} \times n_{\text{Kinetics}}) - 1$, positive when ETRI tends to have larger raw values.

```
[22]: def dataset_comparisons(data):
    all_families = []
    for metric in METRICS:
        rows = []
        for method in METHODS:
            subset = data[data['metric'].eq(metric) & data['method'].eq(method)]
            etri = subset.loc[subset['dataset'].eq(DATASETS[0]), 'value'].
↳dropna().to_numpy()
            kinetics = subset.loc[subset['dataset'].eq(DATASETS[1]), 'value'].
↳dropna().to_numpy()
            if len(etri) and len(kinetics):
                result = stats.mannwhitneyu(etri, kinetics,
↳alternative='two-sided', method='auto')
                u_stat, p_raw = float(result.statistic), float(result.pvalue)
                effect = 2 * u_stat / (len(etri) * len(kinetics)) - 1
            else:
                u_stat, p_raw, effect = np.nan, np.nan, np.nan
            rows.append({
                'Metric': metric, 'Method': DISPLAY_NAMES[method],
                'N ETRI': len(etri), 'N Kinetics': len(kinetics),
```

```

        'Mann-Whitney U': u_stat, 'p raw': p_raw,
        'Rank-biserial effect (ETRI-Kinetics)': effect,
        'Median ETRI': np.median(etri) if len(etri) else np.nan,
        'Median Kinetics': np.median(kinetics) if len(kinetics) else np.
↪nan,
    })
    family = pd.DataFrame(rows)
    valid = family['p raw'].notna()
    family['p Holm'] = np.nan
    family.loc[valid, 'p Holm'] = holm_adjust(family.loc[valid, 'p raw'])
    family['Significant'] = family['p Holm'] < ALPHA
    all_families.append(family)
    return pd.concat(all_families, ignore_index=True)

dataset_comparison_df = dataset_comparisons(metric_long)
dataset_comparison_df.to_csv(RESULTS_DIR / 'dataset_mannwhitney_holm.csv',
↪index=False)
display(dataset_comparison_df.round(5))

```

	Metric	Method	N ETRI	N Kinetics	Mann-Whitney U	\
0	deletion_auc	3D-Kernel-SHAP	30	30	313.0	
1	deletion_auc	3D-LIME	30	30	337.0	
2	deletion_auc	3D-RISE	30	30	276.0	
3	deletion_auc	3D-Sampled-Occl-Sens	30	30	471.0	
4	deletion_auc	LV-LOCO	30	30	365.0	
..	...	...	...	...	...	
94	sensitivity_n	LV-LOCO	30	30	583.0	
95	sensitivity_n	LV-Univ-Pred	30	30	591.0	
96	sensitivity_n	SaliencyTubes	30	30	415.0	
97	sensitivity_n	AOSA	30	30	379.0	
98	sensitivity_n	GradCAM	30	30	356.0	

	p raw	Rank-biserial effect (ETRI-Kinetics)	Median ETRI	\
0	0.04358	-0.30444	0.08995	
1	0.09626	-0.25111	0.07603	
2	0.01031	-0.38667	0.09893	
3	0.76183	0.04667	0.31645	
4	0.21156	-0.18889	0.15779	
..	...	...	...	
94	0.05012	0.29556	0.51696	
95	0.03778	0.31333	0.09874	
96	0.61001	-0.07778	0.17503	
97	0.29727	-0.15778	0.14557	
98	0.16687	-0.20889	0.16865	

	Median Kinetics	p Holm	Significant
0	0.15940	0.30508	False

1	0.12289	0.57758	False
2	0.17600	0.08252	False
3	0.24076	1.00000	False
4	0.22414	1.00000	False
..	...	...	...
94	0.41022	0.40096	False
95	-0.00696	0.34004	False
96	0.28477	1.00000	False
97	0.20990	1.00000	False
98	0.34175	0.83433	False

[99 rows x 11 columns]


1.10 11. Visualizations


Each discovered metric gets one compact three-panel boxplot (overall, ETRI, Kinetics). Within a panel, methods are ordered by median in the configured performance direction. Boxplots show distributions, pairing is handled in the tests rather than drawn as 60×9 connecting lines, and no visual ordering should be read as proof of significance.





```
[23]: def plot_metric_by_scope(data, metric):
    fig, axes = plt.subplots(1, 3, figsize=(18, 5))
    higher = METRIC_CONFIG[metric]['higher_is_better']
    for ax, scope in zip(axes, SCOPES):
        subset = data[data['metric'].eq(metric)]
        if scope != 'Overall':
            subset = subset[subset['dataset'].eq(scope)]
        order = (subset.groupby('method')['value'].median()
                 .sort_values(ascending=not higher).index.tolist())
        values = [subset.loc[subset['method'].eq(method), 'value'].dropna().
        ↪to_numpy() for method in order]
        ax.boxplot(values, tick_labels=[DISPLAY_NAMES[method] for method in
        ↪order],
                    showmeans=True, meanprops={'marker': 'o', 'markerfacecolor':
        ↪'white',
                                                'markeredgecolor': 'black',
        ↪'markersize': 4})
        ax.set_title(scope)
        ax.tick_params(axis='x', rotation=55)
        for label in ax.get_xticklabels():
            label.set_ha('right')
        ax.set_ylabel(metric if ax is axes[0] else '')
        direction = 'higher is better' if higher else 'lower is better'
        fig.suptitle(f'{metric} ({direction})', y=1.02, fontsize=13)
        fig.tight_layout()
        fig.savefig(FIGURES_DIR / f'{metric}_boxplots.png', dpi=160,
        ↪bbox_inches='tight')
```

```
plt.show()
```

```
for metric in METRICS:
    plot_metric_by_scope(metric_long, metric)
```







1.11 12. Summary of Statistical Findings

The following text and compact tables are generated from the computed results. They deliberately separate numerical winners from statistically supported pairwise differences and report taxonomy metrics that could not enter a repeated-video analysis.

```
[24]: removed_total = int(global_tests_df['Rows removed'].sum())
significant_global = int(global_tests_df['Significant'].sum())
significant_posthoc = int(posthoc_df['Significant'].sum()) if not posthoc_df.
    empty else 0
significant_dataset = int(dataset_comparison_df['Significant'].sum())
overall_best = ranking_df[ranking_df['Dataset'].eq('Overall') &
    ranking_df['Best']]
best_text = '; '.join(
    f"{row.Metric}: {row.Method}" for row in overall_best[['Metric', 'Method']].
    itertuples(index=False)
)
display(Markdown(
    f"**Computed summary (alpha = {ALPHA}):** {significant_global} of
    {len(global_tests_df)} "
    f"Friedman tests were significant. Their post-hoc families contained "
    f"{significant_posthoc} Holm-significant pairwise comparisons.
    Complete-case filtering removed "
    f"{removed_total} block rows across all metric/scope tests (a row can be
    counted in more than one scope). "
    f"The independent dataset analysis found {significant_dataset}
    Holm-significant method-level comparisons."
))
display(Markdown("**Numerical overall leaders (not significance claims):** " +
    best_text))
unavailable = metric_availability_df.loc[
    ~metric_availability_df['Included in repeated-video analysis'], ['Metric',
    Reason']
]
if not unavailable.empty:
    display(Markdown("**Unavailable/excluded taxonomy metrics:** " + '; '.join(
        f"{row.Metric} - {row.Reason}" for row in unavailable.
        itertuples(index=False)
    )))

final_global_tests = global_tests_df[[
    'Metric', 'Dataset', 'N', 'Friedman statistic', 'p', 'Kendall W',
    Significant'
]]
final_posthoc = posthoc_df[[
    'Metric', 'Dataset', 'Method A', 'Method B', 'p raw', 'p Holm',
    'Rank-biserial effect', 'Significant'
```

```

]]
final_ranking = ranking_df[[
    'Metric', 'Dataset', 'Rank', 'Method', 'Mean', 'Median'
]]
display(Markdown('### Global tests'))
display(final_global_tests.round(5))
display(Markdown('### Post-hoc tests'))
display(final_posthoc.round(5))
display(Markdown('### Rankings'))
display(final_ranking.round(4))

```

Computed summary (alpha = 0.05): 32 of 33 Friedman tests were significant. Their post-hoc families contained 565 Holm-significant pairwise comparisons. Complete-case filtering removed 0 block rows across all metric/scope tests (a row can be counted in more than one scope). The independent dataset analysis found 30 Holm-significant method-level comparisons.

Numerical overall leaders (not significance claims): avg_drop: 3D-RISE; bbox_energy_ratio: GradCAM; deletion_auc: 3D-LIME; insertion_auc: 3D-LIME; iou50: LV-LOCO; iou_value: LV-LOCO; pointing_game: 3D-LIME; runtime_seconds: GradCAM; runtime_seconds: SaliencyTubes; sensitivity_n: LV-LOCO; sparsity_80: GradCAM; temporal_consistency: SaliencyTubes

Unavailable/excluded taxonomy metrics: roar — Excluded: results are dataset/method/fraction-level, not per-video repeated measures.

1.11.1 Global tests

	Metric	Dataset	N	Friedman statistic	p	\
0	deletion_auc	Overall	60	136.43556	0.00000	
1	deletion_auc	EtriActivity3D	30	83.44000	0.00000	
2	deletion_auc	Kinetics400	30	69.07556	0.00000	
3	insertion_auc	Overall	60	115.50667	0.00000	
4	insertion_auc	EtriActivity3D	30	108.56000	0.00000	
5	insertion_auc	Kinetics400	30	79.10222	0.00000	
6	avg_drop	Overall	60	187.90009	0.00000	
7	avg_drop	EtriActivity3D	30	94.29634	0.00000	
8	avg_drop	Kinetics400	30	133.87240	0.00000	
9	pointing_game	Overall	60	35.42246	0.00002	
10	pointing_game	EtriActivity3D	30	26.05025	0.00103	
11	pointing_game	Kinetics400	30	18.92571	0.01526	
12	iou50	Overall	60	25.72727	0.00117	
13	iou50	EtriActivity3D	30	24.70000	0.00175	
14	iou50	Kinetics400	30	9.00000	0.34230	
15	iou_value	Overall	60	42.14934	0.00000	
16	iou_value	EtriActivity3D	30	54.82661	0.00000	
17	iou_value	Kinetics400	30	18.33614	0.01884	
18	bbox_energy_ratio	Overall	60	149.27733	0.00000	
19	bbox_energy_ratio	EtriActivity3D	30	115.20000	0.00000	
20	bbox_energy_ratio	Kinetics400	30	53.74708	0.00000	

21	temporal_consistency	Overall	60	324.61333	0.00000
22	temporal_consistency	EtriActivity3D	30	170.67556	0.00000
23	temporal_consistency	Kinetics400	30	180.71111	0.00000
24	sparsity_80	Overall	60	194.64444	0.00000
25	sparsity_80	EtriActivity3D	30	96.04444	0.00000
26	sparsity_80	Kinetics400	30	129.11111	0.00000
27	runtime_seconds	Overall	60	475.00728	0.00000
28	runtime_seconds	EtriActivity3D	30	238.23417	0.00000
29	runtime_seconds	Kinetics400	30	236.83585	0.00000
30	sensitivity_n	Overall	60	129.91556	0.00000
31	sensitivity_n	EtriActivity3D	30	82.59556	0.00000
32	sensitivity_n	Kinetics400	30	68.48000	0.00000

	Kendall W	Significant
0	0.28424	True
1	0.34767	True
2	0.28781	True
3	0.24064	True
4	0.45233	True
5	0.32959	True
6	0.39146	True
7	0.39290	True
8	0.55780	True
9	0.07380	True
10	0.10854	True
11	0.07886	True
12	0.05360	True
13	0.10292	True
14	0.03750	False
15	0.08781	True
16	0.22844	True
17	0.07640	True
18	0.31099	True
19	0.48000	True
20	0.22395	True
21	0.67628	True
22	0.71115	True
23	0.75296	True
24	0.40551	True
25	0.40019	True
26	0.53796	True
27	0.98960	True
28	0.99264	True
29	0.98682	True
30	0.27066	True
31	0.34415	True
32	0.28533	True

1.11.2 Post-hoc tests

	Metric	Dataset	Method A	Method B \
0	deletion_auc	Overall	3D-Kernel-SHAP	3D-LIME
1	deletion_auc	Overall	3D-Kernel-SHAP	3D-RISE
2	deletion_auc	Overall	3D-Kernel-SHAP	3D-Sampled-Occl-Sens
3	deletion_auc	Overall	3D-Kernel-SHAP	LV-LOCO
4	deletion_auc	Overall	3D-Kernel-SHAP	LV-Univ-Pred
...	...	...	...	...
1147	sensitivity_n	Kinetics400	LV-Univ-Pred	AOSA
1148	sensitivity_n	Kinetics400	LV-Univ-Pred	GradCAM
1149	sensitivity_n	Kinetics400	SaliencyTubes	AOSA
1150	sensitivity_n	Kinetics400	SaliencyTubes	GradCAM
1151	sensitivity_n	Kinetics400	AOSA	GradCAM

	p raw	p Holm	Rank-biserial effect	Significant
0	0.00001	0.00025	0.65355	True
1	0.59099	1.00000	0.07978	False
2	0.00000	0.00004	-0.71694	True
3	0.00001	0.00017	-0.66885	True
4	0.00000	0.00000	-0.85027	True
...	...	...	...	...
1147	0.00067	0.01600	-0.68172	True
1148	0.00000	0.00008	-0.88387	True
1149	0.29884	1.00000	0.22151	False
1150	0.00051	0.01314	-0.69462	True
1151	0.01283	0.24384	-0.51398	False

[1152 rows x 8 columns]

1.11.3 Rankings

	Metric	Dataset	Rank	Method \
0	avg_drop	EtriActivity3D	1.0	3D-RISE
1	avg_drop	EtriActivity3D	2.0	3D-LIME
2	avg_drop	EtriActivity3D	3.0	SaliencyTubes
3	avg_drop	EtriActivity3D	4.0	AOSA
4	avg_drop	EtriActivity3D	5.0	3D-Kernel-SHAP
..	...	...	...	...
292	temporal_consistency	Overall	5.0	3D-Kernel-SHAP
293	temporal_consistency	Overall	6.0	AOSA
294	temporal_consistency	Overall	7.0	LV-LOCO
295	temporal_consistency	Overall	8.0	3D-RISE
296	temporal_consistency	Overall	9.0	3D-Sampled-Occl-Sens
	Mean	Median		
0	2.3191	0.0001		
1	16.3615	0.6543		

2	20.1009	4.8027
3	20.9533	0.4261
4	22.2154	17.1348
..	...	...
292	0.9350	0.9512
293	0.9312	0.9504
294	0.9227	0.9258
295	0.4457	0.4449
296	0.0080	-0.0137

[297 rows x 6 columns]

1.11.4 Reproducible outputs

Tables written to docs/results:

- statistical_metric_availability.csv
- statistical_descriptive_summary.csv
- friedman_global_tests.csv
- wilcoxon_holm_posthoc.csv
- xai_method_rankings.csv
- dataset_mannwhitney_holm.csv

One three-scope boxplot per automatically discovered metric is written to docs/figures/statistical-analysis.